



# CS5220 Advanced Topics in Web Programming

Introduction to MongoDB

Chengyu Sun  
California State University, Los Angeles



# NoSQL

- ◆ Not SQL, Not Only SQL, *Not Relational*
- ◆ A term that describes a class of data storage and manipulation technologies and products that do not follow the RDBMS principles and focus on large datasets, performance, scalability, and agility

# Types of NoSQL Databases

- ◆ Key-Value Stores
- ◆ Column Family Stores
- ◆ Graph Databases
- ◆ Document Databases
  - A *document* in a document database consists of a loosely structured set of key-value pairs.

# A "Document" Example

```
{  
  "first_name": "John",  
  "last_name": "Doe",  
  "age": 20,  
  "address": {  
    "street": "123 Main"  
    "city": "Los Angeles"  
    "state": "CA"  
  }  
}
```

- ◆ It's basically JSON
- ◆ Why is it called a *document*, not a *object*??

# DBMS Popularity

## ◆ DB-Engines.com

- [Ranking](#)
- [Trend](#)

# MongoDB

- ◆ The most popular NoSQL database
  - Document database
  - BSON data types
- ◆ The "M" in MEAN/MERN stack

# MongoDB Installation

## ◆ Install MongoDB Community Edition

- No security by default
- Localhost binding (default since 3.6)

## ◆ Install MongoDB Compass

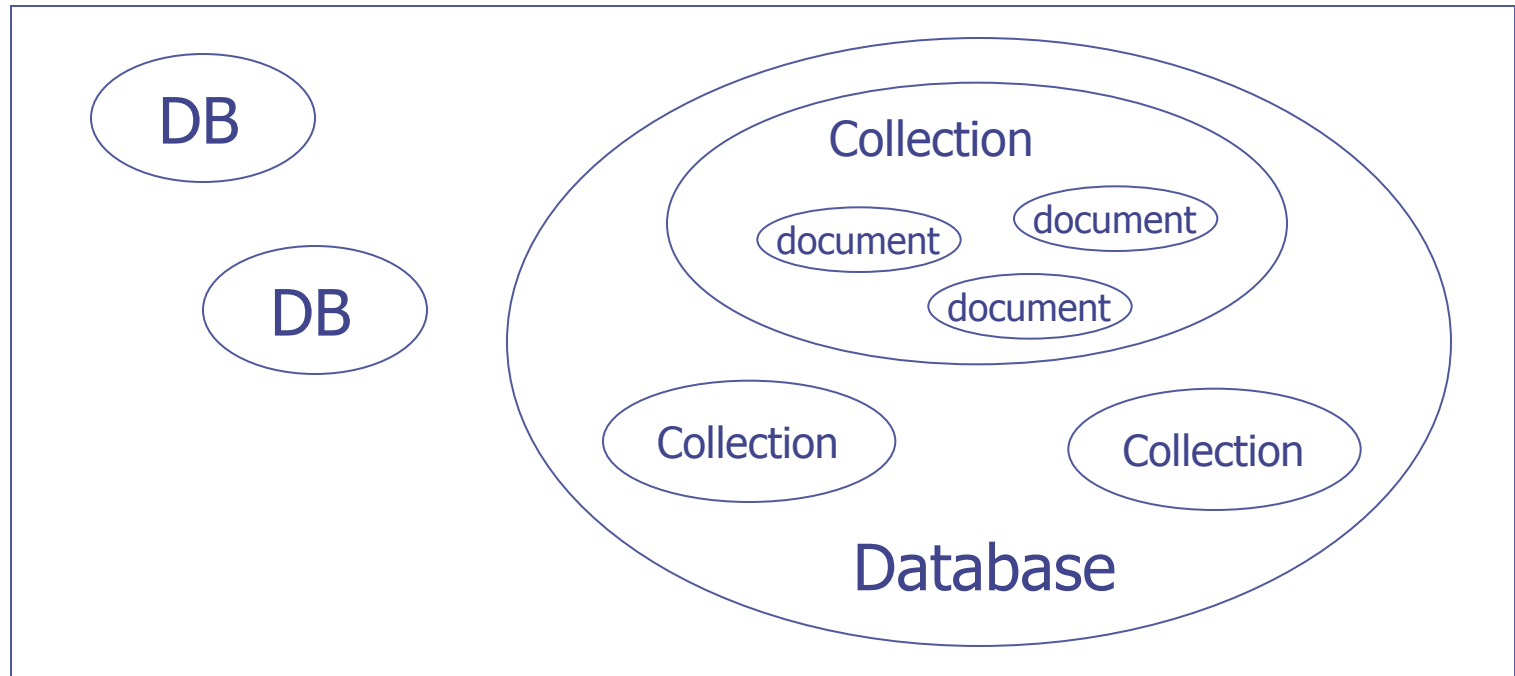
- MongoDB Windows Installer installs Compass by default

# Basic Usage of MongoDB Compass

- ◆ Connection
- ◆ Create databases and collections
- ◆ Document CRUD



# MongoDB Server



- ◆ A **collection** is the equivalent of a table in relational databases
- ◆ Collection does *not* enforce a schema

# Blog Example

- ◆ Users
- ◆ Articles
- ◆ Comments
- ◆ Tags

# Data Modeling

- ◆ Understand MongoDB data types
- ◆ Knowledge of relational modeling still applies
- ◆ To embed or not to embed: that is the question

# Data Modeling (I): Data Types

## ◆ Numbers

- Boolean, Integer, Double, Decimal

## ◆ Text and binary data

- String, Regular Expression, Code, Binary Data

## ◆ Date and timestamp (for internal use)

## ◆ Special types

- Null, Min/Max Key, **ObjectId**

## ◆ Object and Array

# A MongoDB Document ...

```
{
  _id: ObjectId("5a09e956df8d3a91d14628d4"),
  title: "My First Blog Post",
  publishDate: 2018-03-31 20:00:00:00,
  author: {
    name: "John"
    email: "john@localhost"
  },
  tags: ["web development", "mongodb"]
}
```

# ... A MongoDB Document

- ◆ Each document must have a unique `_id` field that serves as the primary key
- ◆ The value of `_id` can be user-assigned (of any type) or auto-generated (of ObjectId type)
- ◆ An `Object` value is also known as an *embedded document* or a *sub-document*

# Data Modeling (II)

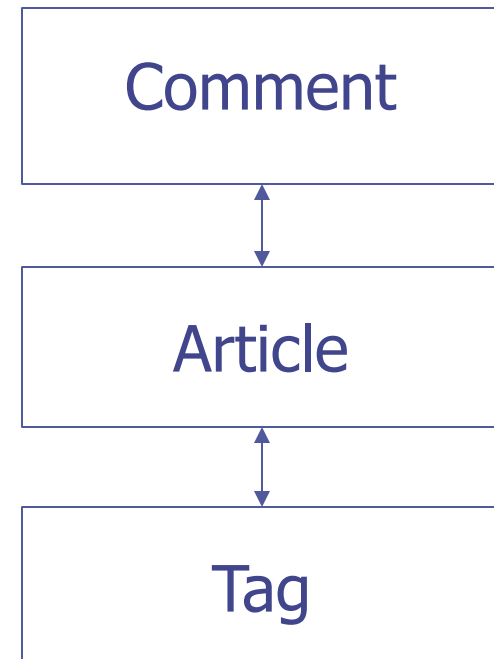
- ◆ All the relational design knowledge still applies
  - Entities and relationships
  - Tables → collections
- ◆ With that said, document databases are not relational databases
  - Doing JOIN is difficult
  - Does have much richer data types

# Data Modeling (III)

- ◆ To embed or not to embed: that's the question



Or





# Embed Or Not Embed

- ◆ Read performance vs data redundancy
  - Basically the same arguments as *normalized* data vs *de-normalized* data in relational design
  - Examples:
    - ◆ Article and article author
    - ◆ Comment and comment author
- ◆ Atomicity requirements

# Need for Transactions ...

◆ Not all operations can be done with a single operation, e.g. transferring money from one bank account to another:

-- 1. Check the balance of account #1  
select balance from accounts where id = 1;

-- 2. Withdraw \$100 from account #1  
update accounts set balance = balance - 100  
where id = 1;

-- 3. Deposit \$100 to account #2  
update accounts set balance = balance + 100  
where id = 2;

# ... Need for Transactions ...

## ◆ Bad things could happen due to concurrent access and/or system failure

-- 1. Check the balance of account #1  
select balance from accounts where id = 1;

*The other owner of the  
joint account withdraws  
all the money in  
account #1*

-- 2. Withdraw \$100 from account #1  
update accounts set balance = balance - 100  
where id = 1;

-- 3. Deposit \$100 to account #2  
update accounts set balance = balance + 100  
where id = 2;

# ... Need for Transactions ...

◆ Bad things could happen due to  
concurrent access and/or system failure

-- 1. Check the balance of account #1

select balance from accounts where id = 1;

*The other owner of the  
joint account checks the  
balances of both  
accounts and notices  
that \$100 is missing*

-- 2. Withdraw \$100 from account #1

update accounts set balance = balance - 100  
where id = 1;

-- 3. Deposit \$100 to account #2

update accounts set balance = balance + 100  
where id = 2;

# ... Need for Transactions

◆ Bad things could happen due to  
concurrent access and/or system failure

-- 1. Check the balance of account #1  
select balance from accounts where id = 1;

-- 2. Withdraw \$100 from account #1  
update accounts set balance = balance - 100  
where id = 1;

*System crash*

-- 3. Deposit \$100 to account #2  
update accounts set balance = balance + 100  
where id = 2;

# Transaction and ACID

- ◆ A transaction in RDBMS is a group of SQL statements treated by the DBMS as *a single unit of work*
- ◆ Transactions in RDBMS are ACID
  - Atomic
  - Consistent
  - Isolated
  - Durable

# Transactions in MongoDB

- ◆ Write operations are atomic on the level of a single document
- ◆ Multi-documents transaction support
  - For replica sets since 4.0 (6/2018)
  - For sharded clusters since 4.2 (8/2019)

# MongoDB Shell

- ◆ `mongo`

- ◆ A command line client that provides an interactive JavaScript interface to MongoDB



# Basic MongoDB Shell Commands

◆ `help`

◆ `show dbs`

◆ `use <db>`

- Switch to database `<db>`
- `<db>` won't be created until some data is inserted into it
- `show collections`
- `db.dropDatabase()`

# MongoDB's Query Language

◆ JavaScript with [MongoDB methods](#)

◆ Some collection methods:

- `db.<collection>.insert()`
- `db.<collection>.find()`
- `db.<collection>.update()`
- `db.<collection>.remove()`
- `db.<collection>.drop()`

# Example of Basic Operations

- ◆ Create a database `test1`
- ◆ Create two documents `John` and `Jane`
- ◆ Save the two documents to a collection  
`users`
- ◆ List the documents in the collection

# Mongo Shell Script

## ◆ Example: `test2.js`

- `connect(<url>)`
- `print()`
- `printjson()`
- `cursor`

## ◆ Run Mongo shell script

- `mongo <script>, or`
- `load("<script>")` inside Mongo shell

# About Mongo Shell Script

- ◆ Don't use shell helper commands like `show dbs` as they are not valid JavaScript
- ◆ Don't use fancy JavaScript syntax as the may not be supported by Mongo shell's JavaScript interpreter

# Using find()

◆ `find( [query], [projection] )`

- [Query tutorials](#)
- [Query operators](#)

◆ Examples using the Blogs database

# Queries (I) Basic Conditions and Projections

- ◆ List all users
- ◆ List the first name of all users
  - Without `_id`
- ◆ Find the users whose last name is Doe
  - Using [comparison operators](#)

# Queries (II) Logical Operators

- ◆ Find the users whose first name is John and last name is Doe
  - Implicit and explicit `$and`
- ◆ Find the users whose first name is John or last name is Doe
- ◆ Find the users whose first name is John or (the first name is Jane and the last name is Doe)



# Queries (III) Arrays and Subdocuments

- ◆ Find the articles whose tags contain "NoSQL"
  - Using `$all`
- ◆ Find the articles John Doe has commented on

# Queries (IV) Join

- ◆ List the articles with their authors (i.e. not just author id)
  - `db.<collection>.aggregate()`
  - [\\$lookup](#)
- ◆ List the article authors
  - [\\$project](#)

# Update and Delete

db.<collection>.update( query, update, options )  
db.<collection>.remove( query, update, options )

- ◆ Change John Doe's name to Tom Smith
  - \$set
  - Other update operators
- ◆ Delete the article "Using MongoDB"
- ◆ Add a tag "Tutorial" to the article "Programming Node.js"
  - \$push
- ◆ Delete the comments made by John Doe
  - \$pull

# Indexing

- ◆ Indexes are crucial for performance just like in RDBMS
- ◆ db.<collection>.createIndex(<keys>, [options])
  - keys: {field: -1|1, field2: -1|1, ...}
  - options: unique, name, ...

# Index Examples

Ascending order



```
db.users.createIndex({lastName: 1});
```

```
db.users.createIndex({email: 1}, {  
  unique: true,  
  name: 'EmailIndex'  
});
```

```
db.articles.createIndex({tags: 1});
```



Index on array field

# MongoDB with Node.js

- ◆ MongoDB Node.js driver

- ◆ Mongoose

- Model classes
  - ◆ Validation
- DAO methods

# Using MongoDB Driver

◆ `npm install mongodb`

◆ Understand the API

- Connection string format
- MongoClient, Db, Collection, Cursor

# About Using MongoDB Driver

- ◆ Difference between Node.js code and MongoDB shell script
- ◆ MongoClient maintains its own connection pool – reuse the same client/db/collection as much as possible before closing it



# Support for Other Programming Languages

- ◆ Drivers for various server-side programming language – <https://docs.mongodb.org/ecosystem/drivers/>

# NoSQL vs Relational ...

- ◆ NoSQL databases are designed to be easier to scale horizontally
- ◆ Document databases make data modeling easier and data access more efficient for certain applications
- ◆ Using one language (i.e. JavaScript) for everything is appealing
- ◆ But ...

# ... NoSQL vs Relational

- ◆ Giving up ACID, SQL, and tried-and-true reliability of RDMBS is no small sacrifice
- ◆ Scalability and performance of RDBMS have continued to improve
- ◆ Understanding application requirements and data modeling is important for both
- ◆ Use the right database for the right job

# Readings

- ◆ MongoDB Manual
- ◆ MongoDB Node.js Driver
- ◆ Why SQL is beating NoSQL